

Joe Benton

+1 (510) 646 3943 | joe.benton127@gmail.com | 2634 Dana St, Berkeley, CA 94704

EDUCATION

D.Phil. in Statistics

2021 – 2023

Department of Statistics, University of Oxford

Thesis title: *Generative Modelling: Theory and Applications*, supervised by Arnaud Doucet and George Deligiannidis

B.A. with M.Math. in Mathematics

2017 – 2021

Trinity College, University of Cambridge

Graduated with Distinction (6th out of c. 250 in year)

Thesis title: *Activated Random Walks*, supervised by Perla Sousi

PROFESSIONAL EXPERIENCE

Member of Technical Staff | *Anthropic*

Fall 2023 – Present

Developed and implemented adversarial testing methods for state-of-the-art large language models. Contributed to plans for producing positive safety cases and aligning ASL-4 level systems as part of the Alignment Science team.

Technical Staff | *UK Frontier AI Taskforce*

Fall 2023

Part of the founding technical team of the UK Frontier AI Taskforce (now UK AI Safety Institute). Produced technical demonstrations of AI risk for presentation in front of world leaders at the first AI Safety Summit.

Machine Learning Researcher | *Redwood Research*

Winter 2022 – 2023

Studied and developed causality and fine-tuning based interpretability methods, with applications to mechanistic anomaly detection. Supervised a month-long intern project aiming to automate interpretability techniques.

Research Intern | *Alignment Research Center*

Spring 2023

Worked on formalizing heuristic arguments, finding efficient heuristic estimators for sparse covariance propagation.

Research Assistant | *Center for Human-Compatible Artificial Intelligence, UC Berkeley*

Summer 2021

Built on the PAIRED algorithm for unsupervised environment design to incorporate human feedback with the aim of speeding up and simplifying the training process. Supervised by Michael Dennis.

VOLUNTEERING

Research Mentor | *Supervised Program for Alignment Research*

Spring – Autumn 2022

Mentored a student research project on decoding sparse feature representations for neural network interpretability.

Trustee, Cofounder | *Raise: A Celebration of Giving*

2018 – Present

Trustee and co-founder of Raise, a student charity initiative raising over £460,000 for the [Against Malaria Foundation](#).

Strategic Advisor | *AI Safety Hub*

2022 – 2023

Advised an AI safety outreach and mentoring organization, and managed their [AI Safety Fundamentals](#) program.

AWARDS

International Mathematics Olympiad (1 Gold – 7th out of 615, 3 Silver)

2014 – 2017

International Olympiad in Informatics (1 Gold – 6th out of 304, 1 Silver, 1 Bronze)

2015 – 2017

Romanian Masters in Mathematics (3 Gold – Best record of any competitor)

2015 – 2017

INVITED TALKS

University of Warwick Algorithms and Computationally Intensive Inference Seminars, 2023

Royal Statistical Society International Conference, Probabilistic and Statistical Aspects of Machine Learning Discussion Meeting, 2023

REVIEWING

JMLR, TMLR, JRSS-B, ICML 2023 Workshop Frontiers4LCD, Cooperative AI Foundation

PUBLICATIONS

From Denoising Diffusions to Denoising Markov Models. **Joe Benton**, Yuyang Shi, Valentin De Bortoli, George Deligiannidis, Arnaud Doucet. *Journal of the Royal Statistical Society, Series B*, 2024. (To appear)

Measuring Feature Sparsity in Language Models. Mingyang Deng, Lucas Tao, **Joe Benton**. *NeurIPS 2023 Workshop on Socially Responsible Language Modelling Research*, 2023

Linear Convergence Bounds for Diffusion Models via Stochastic Localization. **Joe Benton**, Arnaud Doucet, George Deligiannidis. *arXiv preprint, arXiv:2308.03686*, 2023

Error Bounds for Flow Matching Methods. **Joe Benton**, George Deligiannidis, Arnaud Doucet. *arXiv preprint, arXiv:2305.16860*, 2023

Alpha-divergence Variational Inference Meets Importance Weighted Auto-Encoders: Methodology and Asymptotics. Kamélia Daudel, **Joe Benton**^{*}, Yuyang Shi^{*}, Arnaud Doucet. *Journal of Machine Learning Research*, 24(243):1–83, 2023

Polysemanticity and Capacity in Neural Networks. Adam Scherlis, Kshitij Sachan, Adam S. Jermyn, **Joe Benton**, Buck Shlegeris. *arXiv preprint, arXiv:2210.01892*, 2022

A Continuous Time Framework for Discrete Denoising Models. Andrew Campbell, **Joe Benton**, Valentin De Bortoli, Tom Rainforth, George Deligiannidis, Arnaud Doucet. *Advances in Neural Information Processing Systems*, 2022