

Joe Benton

EDUCATION

D.Phil. in Statistics

October 2021 – August 2023

Department of Statistics, University of Oxford

Thesis title: *Generative Modelling: Theory and Applications*, supervised by Arnaud Doucet and George Deligiannidis

B.A. with M.Math. in Mathematics

September 2017 – June 2021

Trinity College, University of Cambridge

Graduated with Distinction (6th out of c. 250 in year)

Thesis title: *Activated Random Walks*, supervised by Perla Sousi

PROFESSIONAL EXPERIENCE

Member of Technical Staff | *Anthropic*

November 2023 – Present

Currently thinking about how to make safety cases for transformative AI systems. Lead contributor to design and production of ASL-4 misalignment evaluations. Previously worked on white-box adversarial testing of LLMs.

Technical Staff | *UK Frontier AI Taskforce*

August 2023 – November 2023

Part of the founding technical team of the UK Frontier AI Taskforce (now UK AI Safety Institute). Produced technical demonstrations of AI risk for presentation in front of world leaders at the first AI Safety Summit.

Machine Learning Researcher | *Redwood Research*

December 2022 – March 2023

Developed and studied causality and fine-tuning based interpretability methods, with applications to mechanistic anomaly detection. Supervised an internship project aiming to automate interpretability techniques.

Research Assistant | *Center for Human-Compatible AI, UC Berkeley*

June 2021 – September 2021

Extended PAIRED algorithm for unsupervised environment design to incorporate human feedback to speeding up and simplify model training. Supervised by Michael Dennis.

VOLUNTEERING

Research Mentor | *ML Alignment & Theory Scholars*

January 2025 – Present

Mentored three students on research projects related to AI control and chain-of-thought faithfulness.

Research Mentor | *Supervised Program for Alignment Research*

March 2023 – September 2023

Mentored a student research project on decoding sparse feature representations for neural network interpretability.

Trustee, Cofounder | *Raise: A Celebration of Giving*

February 2018 – Present

Trustee and co-founder of Raise, a student charity initiative raising over £460,000 for the [Against Malaria Foundation](#).

AWARDS

International Mathematics Olympiad (1 Gold – 7th out of 615, 3 Silver)

2014 – 2017

International Olympiad in Informatics (1 Gold – 6th out of 304, 1 Silver, 1 Bronze)

2015 – 2017

Romanian Masters in Mathematics (3 Gold – Best record of any competitor)

2015 – 2017

INVITED TALKS

University of Warwick Algorithms and Computationally Intensive Inference Seminars, 2023

Royal Statistical Society International Conference, Probabilistic and Statistical Aspects of Machine Learning Discussion Meeting, 2023

SERVICE

Program Chair, AAAI 2025 Alignment Track

Reviewed for ICLR 2025, NeurIPS 2024, JMLR, TMLR, JRSS-B, NeurIPS 2023 SoLaR Workshop, NeurIPS 2023

Workshop on Diffusion Models, ICML 2023 Workshop Frontiers4LCD, Cooperative AI Foundation

PUBLICATIONS

- Constitutional Classifiers: Defending against Universal Jailbreaks across Thousands of Hours of Red Teaming.* Mrinank Sharma, Meg Tong, Jesse Mu, Jerry Wei, Jorrit Kruthoff, Scott Goodfriend, Euan Ong, Alwin Peng, Raj Agarwal, Cem Anil, Amanda Askill, Nathan Bailey, **Joe Benton**, Emma Blumke, et al. *arXiv preprint arXiv:2501.18837*, 2025.
- Sabotage Evaluations for Frontier Models.* **Joe Benton**, Misha Wagner, Eric Christiansen, Cem Anil, Ethan Perez, Jai Srivastav, Esin Durmus, Deep Ganguli, Shauna Kravec, Buck Shlegeris et al. *arXiv preprint arXiv:2410.21514*, 2024.
- Many-shot Jailbreaking.* Cem Anil, Esin Durmus, Mrinank Sharma, **Joe Benton**, Sandipan Kundu, Joshua Batson, Nina Rimskey, Meg Tong, Jesse Mu, Daniel Ford et al. *Advances in Neural Information Processing Systems*, 2024.
- When Do Universal Image Jailbreaks Transfer Between Vision-Language Models?.* Rylan Schaeffer, Dan Valentine, Luke Bailey, James Chua, Cristóbal Eyzaguirre, Zane Durante, **Joe Benton**, Brando Miranda, Henry Sleight, John Hughes et al. *NeurIPS 2024 Workshop on Responsibly Building the Next Generation of Multimodal Foundational Models*, 2024.
- From Denoising Diffusions to Denoising Markov Models.* **Joe Benton**, Yuyang Shi, Valentin De Bortoli, George Deligiannidis, Arnaud Doucet. *Journal of the Royal Statistical Society Series B: Statistical Methodology*, 86(2):286–301, 2024.
- Nearly d -Linear Convergence Bounds for Diffusion Models via Stochastic Localization.* **Joe Benton**, Arnaud Doucet, George Deligiannidis. *International Conference on Learning Representations*, 2024.
- Error Bounds for Flow Matching Methods.* **Joe Benton**, George Deligiannidis, Arnaud Doucet. *Transactions on Machine Learning Research*, February 2024.
- Alpha-divergence Variational Inference Meets Importance Weighted Auto-Encoders: Methodology and Asymptotics.* Kamliia Daudel, **Joe Benton***, Yuyang Shi*, Arnaud Doucet. *Journal of Machine Learning Research*, 24(243):1–83, 2023.
- Measuring Feature Sparsity in Language Models.* Mingyang Deng, Lucas Tao, **Joe Benton**. *NeurIPS 2023 Workshop on Socially Responsible Language Modelling Research*, 2023.
- A Continuous Time Framework for Discrete Denoising Models.* Andrew Campbell, **Joe Benton**, Valentin De Bortoli, Tom Rainforth, George Deligiannidis, Arnaud Doucet. *Advances in Neural Information Processing Systems*, 2022.
- Polysemanticity and Capacity in Neural Networks.* Adam Scherlis, Kshitij Sachan, Adam S. Jermyn, **Joe Benton**, Buck Shlegeris. *arXiv preprint, arXiv:2210.01892*, 2022.